

# Parallel Programming

## Basic MPI

Timothy H. Kaiser, Ph.D.  
[tkaiser2@nrel.gov](mailto:tkaiser2@nrel.gov)

# Examples and Slides

Examples:

```
git clone https://github.com/timkphd/examples  
cd examples/mpi
```

Slides:

```
https://github.com/timkphd/slides
```



# Talk Overview

- Background on MPI
- Documentation
- Hello world in MPI
- Basic communications
- Simple send and receive program

# Background on MPI

- MPI - Message Passing Interface
  - Library standard defined by a committee of vendors, implementers, & parallel programmers
  - Used to create **parallel programs** based on message passing
- Portable: one standard, many implementations
- Available on almost all parallel machines in **C and Fortran**
- Over 100 advanced routines but 6 basic
- Also used to create **libraries**.



# Unofficial Languages

- Subsets available for R and Java
- Fairly complete implementation for Python
- Let me know if you are interested in these
- I have not looked at Julia

# Documentation

- MPI home page (contains the library standard)  
<https://www.mpi-forum.org>
- Books
  - "MPI: The Complete Reference" by Snir, Otto, Huss-Lederman, Walker, and Dongarra, MIT Press (also in Postscript and html)
  - "Using MPI" by Gropp, Lusk and Skjellum, MIT Press
- Tutorials
  - many online, just do a search



# MPI Implementations

- Most parallel supercomputer vendors provide optimized implementations
- Intel
- HP/SGI/Cray - mpt
- IBM

# Open Source MPI Implementations

- MPICH:
  - <https://www.mpich.org>
- Mvapich
  - <http://mvapich.cse.ohio-state.edu>
- OpenMPI
  - [www.open-mpi.org](http://www.open-mpi.org)



# Key Concepts of MPI

- Used to create parallel programs based on message passing
- Normally the same program is running on several different processors
- Processors **communicate** using **message** passing
- Typical methodology:

```
start job on n processors
do i=1 to j
  each processor does some calculation
  pass messages between processor
end do
end job
```

# Communicators

- Communicator
  - A collection of processors working on some part of a parallel job
  - Used as a parameter for most MPI calls
  - `MPI_COMM_WORLD` includes all of the processors in your job
  - Processors within a communicator are assigned numbers (ranks) 0 to  $n-1$
  - Can create subsets of `MPI_COMM_WORLD`



# Messages

- Simplest message: an array of data of one type.
- Predefined types correspond to commonly used types in a given language
  - MPI\_REAL (Fortran), MPI\_FLOAT (C)
  - MPI\_DOUBLE\_PRECISION (Fortran), MPI\_DOUBLE (C)
  - MPI\_INTEGER (Fortran), MPI\_INT (C)
- User can define more complex types and send packages.

# Include files

- The MPI include file
  - C: `mpi.h`
  - Fortran: `mpif.h` (a f90 module is a good place for this)
- Defines many constants used within MPI programs
- In C defines the interfaces for the functions
- Compilers know where to find the include files



# Minimal MPI program

- Every MPI program needs these...

- C version

```
/* the mpi include file */
#include <mpi.h>
    int nPES, ierr, iam;
/* Initialize MPI */
    ierr=MPI_Init(&argc, &argv);
/* How many processors (nPES) are there? */
    ierr=MPI_Comm_size(MPI_COMM_WORLD, &nPES);
/* What processor am I (what is my rank)? */
    ierr=MPI_Comm_rank(MPI_COMM_WORLD, &iam);
...
    ierr=MPI_Finalize();
```

In C MPI routines are functions and return an error value

# Minimal MPI program

- Every MPI program needs these...
  - Fortran version

```
! MPI include file
  include 'mpif.h'
! The mpi module can be used for Fortran 90 instead of mpif.h
!   use mpi
  integer nPEs, ierr, iam
! Initialize MPI
  call MPI_Init(ierr)
! How many processors (nPEs) are there?
  call MPI_Comm_size(MPI_COMM_WORLD, nPEs, ierr)
! What processor am I (what is my rank)?
  call MPI_Comm_rank(MPI_COMM_WORLD, iam, ierr)
  ...
  call MPI_Finalize(ierr)
```

In Fortran, MPI routines are subroutines, and last parameter is an error value



# Exercise I : Hello World

- Write a parallel “hello world” program
  - Initialize MPI
  - Have each processor print out “Hello, World” and its processor number (rank)
  - Quit MPI

# Fortran and C examples

```
el2:mpi> cat helloc.c
#include <stdio.h>
#include <stdlib.h>
#include <mpi.h>
#include <math.h>

/*****
This is a simple hello world program. Each processor prints
name, rank, and total run size.
*****/
int main(int argc, char **argv)
{
    int myid,numprocs,resultlen;
    char myname[MPI_MAX_PROCESSOR_NAME] ;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);

    MPI_Get_processor_name(myname,&resultlen);

    printf("Hello from %s %d %d\n",myname,myid,numprocs);
    MPI_Finalize();
}
```



# Fortran and C examples

```
el2:mpi> cat hellof.f90
```

```
!*****  
! This is a simple hello world program. Each processor  
! prints out its name, rank and number of processors  
! in the current MPI run.  
!*****  
    program hello  
    include "mpif.h"  
    integer myid,numprocs,ierr,nlength  
    character (len=MPI_MAX_PROCESSOR_NAME):: myname  
    call MPI_INIT( ierr )  
    call MPI_COMM_RANK( MPI_COMM_WORLD, myid, ierr )  
    call MPI_COMM_SIZE( MPI_COMM_WORLD, numprocs, ierr )  
  
    call MPI_Get_processor_name(myname,nlength,ierr)  
  
    write (*,*) "Hello from ",trim(myname)," # ",myid," of ",numprocs  
    call MPI_FINALIZE(ierr)  
    stop  
end  
el2:mpi>  
el2:mpi>
```

# Compiling

- Most MPI compilers are actually just scripts that call underlying Fortran or C compilers
- Load module that points to your compilers
- For Intel compilers use
  - `module load intel-mpi`
  - `mpiifort`, `mpiicc`, `mpicpc` - Intel backend compilers
  - `mpif90` `mpicc`, `mpicxx` - gcc/gfortran backend compilers
- For HP's version of MPI
  - `module load mpt gcc/8.4.0`
  - `mpif77` `mpif90`
  - `mpicc` `mpiCC`



# Compiling IntelMPI and MPT

```
el2:mpi> ml intel-mpi
el2:mpi>
el2:mpi> mpiicc helloc.c -o helloc
el2:mpi> mpiifort hellof.f90 -o hellof
el2:mpi> ##### or #####
el2:mpi> mpicc helloc.c -o helloc
el2:mpi> mpif90 hellof.f90 -o hellof
```

Or

```
el2:mpi> module load mpt gcc/8.4.0
el2:mpi>
el2:mpi> mpicc helloc.c -o hello
el2:mpi> mpif90 hellof.f90 -o hellof
el2:mpi>
```

# Compiling OpenMPI

- Most MPI compilers are actually just scripts that call underlying Fortran or C compilers
- Load module that points to your compilers
- Openmpi
  - `module load openmpi`
  - `mpif77 mpif90`
  - `mpicc mpiCC mpic++ mpicxx`
- This will give you gcc/gfortran as the backend compilers
- To get Intel backend compilers
  - `export OMPI_FC=ifort`
  - `export OMPI_CC=icc`
  - `export OMPI_CXX=icpc`



# makefile

```
el1:mpi> ml intel-mpi
```

```
el1:mpi> make
```

```
mpicc      c_ex00.c -o c_ex00
```

```
mpicc      c_ex01.c -o c_ex01
```

```
mpicc      c_ex02.c -o c_ex02
```

```
mpicc      c_ex03.c -o c_ex03
```

```
mpicc      c_ex04.c -o c_ex04
```

```
mpicc      c_ex05.c -o c_ex05
```

```
mpicc      c_ex06.c -o c_ex06
```

```
mpicc      c_ex07.c -o c_ex07
```

```
mpicc      c_ex08.c -o c_ex08
```

```
mpicc      c_ex09.c -o c_ex09
```

```
mpicc      c_ex10.c -o c_ex10
```

```
mpicc      c_ex11.c -o c_ex11
```

```
mpicc      c_ex12.c -o c_ex12
```

```
mpicc      c_ex13.c -o c_ex13
```

```
c_ex13.c: In function 'main':
```

```
c_ex13.c:34:23: warning: cast to pointer from integer of different size [-Wint-to-pointer-cast]
```

```
    mpi_err = MPI_Gather((void*)mysize,1,MPI_INT,
```

```
mpicc      helloc.c -o helloc
```

```
mpif90     f_ex00.f90 -o f_ex00
```

```
mpif90     f_ex01.f90 -o f_ex01
```

```
mpif90     f_ex02.f90 -o f_ex02
```

```
mpif90     f_ex03.f90 -o f_ex03
```

```
mpif90     f_ex04.f90 -o f_ex04
```

```
mpif90     f_ex05.f90 -o f_ex05
```

```
mpif90     f_ex06.f90 -o f_ex06
```

```
mpif90     f_ex07.f90 -o f_ex07
```

```
mpif90     f_ex08.f90 -o f_ex08
```

```
mpif90     f_ex09.f90 -o f_ex09
```

```
mpif90     f_ex10.f90 -o f_ex10
```

```
mpif90     f_ex11.f90 -o f_ex11
```

```
mpif90     f_ex12.f90 -o f_ex12
```

```
mpif90     f_ex13.f90 -o f_ex13
```

```
mpif90     hellof.f90 -o hellof
```

```
el1:mpi>
```

# Running

- Most often you will use a batch system
- Write a batch script file.
- Use the command **~~mpiexec~~** or **~~mpirun~~ srun** to actually start the program. You must tell the system how many copies to run
- On some systems you must tell where to run the program



# A very simple Slurm Script - "sim"

```
#!/bin/bash
#SBATCH --job-name="flow"
#SBATCH --nodes=1
#SBATCH --export=ALL
#SBATCH --oversubscribe
#SBATCH --time=00:10:00
#SBATCH --partition=debug
#SBATCH --account=hpcapps
```

```
module purge
ml intel-mpi
```

```
srun -n 4 ./helloc
```

The run...

```
el1:mpi> sbatch --account=hpcapps sim
Submitted batch job 8182395
el1:mpi> cat slurm-8182395.out
Hello from r1i7n35 #      2 of      4
Hello from r1i7n35 #      0 of      4
Hello from r1i7n35 #      1 of      4
Hello from r1i7n35 #      3 of      4
el1:mpi>
```

# A useful example - "doeagle"


```
#!/bin/bash
#SBATCH --job-name="flow"
#SBATCH --nodes=1
#SBATCH --export=ALL
#SBATCH --oversubscribe
#SBATCH --time=00:10:00
#SBATCH --partition=debug
#SBATCH --account=hpcapps
```

```
module purge
ml intel-mpi
```

```
if [ -z ${EXE+x} ]; then export EXE=c_ex00 ; fi
```

```
echo "running " $EXE
srun -n 4 ./$EXE
```

You can set the  
program to run with  
and export before  
you submit the run



```
el1:mpi> export EXE=f_ex02
el1:mpi> sbatch --account=hpcapps doeagle
Submitted batch job 8182393
el1:mpi> cat *8182393*
running f_ex02
Hello from fortran process:
Hello from fortran process:
Hello from fortran process:
Hello from fortran process:
getting 1
i= 100
el1:mpi>
```

The run...

3	Numprocs is	4
1	Numprocs is	4
2	Numprocs is	4
0	Numprocs is	4



# Slight Digression

- One way you can understand what is going on in a MPI application is to look at it running in a debugger
- We have ARM/ddt
- We also have ARM/map and ARM/Profile for profiling
- Let me know if interested

# ARM ddt

Arm DDT - Arm Forge 21.1.2

Current Group: All Focus on current: ☒ Group ☐ Process ☐ Thread ☐ Step Threads Together

All 0 1 2 3

Create Group

Project Files

Search (%K)

Application Code

Sources

c\_ex00.c

main(arg)

External Code

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <mpi.h>
4 #include <math.h>
5
6 /*****
7  This is a simple hello world program. Each processor prints out
8  it's rank and the size of the current MPI run (Total number of
9  processors).
10 *****/
11 int main(argc,argv)
12 int argc;
13 char *argv[];
14 {
15     int myid, numprocs;
16
17     MPI_Init(&argc,&argv);
18     MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
19     MPI_Comm_rank(MPI_COMM_WORLD,&myid);
20
21
22     /* print out my rank and this run's PE size */
23     printf("Hello from task %d of %d\n",myid,numprocs);
24
25     MPI_Finalize();
26 }
27
28
29
```

Locals Current Line(s) Current Stack

Current Line(s)

Name	Value
numprocs	0

Breakp... Watchp... **Stacks** Tracep... Tracepoint O...

Stacks

Processes	Threads	Function
4	4	main (c_ex00.c:18)
4	4	ucs_async_thread_func

Evaluate

Name	Value
------	-------

Ready Connected to: (via tunnel) el1:4201 -> el1



# Basic Communication

- Data values are transferred from one processor to another
  - One processor sends the data
  - Another receives the data
- Synchronous
  - Call does not return until the message is sent or received
- Asynchronous
  - Call indicates a start of send or receive, and another call is made to determine if finished

# Synchronous Send

- C
  - `MPI_Send(&buffer, count ,datatype, destination, tag,communicator);`
- Fortran
  - `Call MPI_Send(buffer, count, datatype, destination,tag,communicator, ierr)`
  - Call blocks until message on the way



**Call MPI\_Send(buffer, count, datatype,  
destination, tag, communicator, ierr)**

- **Buffer**: The data array to be sent
- **Count** : Length of data array (in elements, 1 for scalars)
- **Datatype** : Type of data, for example : MPI\_DOUBLE\_PRECISION, MPI\_INT, etc
- **Destination** : Destination processor number (within given communicator)
- **Tag** : Message type (arbitrary integer)
- **Communicator** : Your set of processors
- **Ierr** : Error return (Fortran only)

# Synchronous Receive

- C
  - `MPI_Recv(&buffer, count, datatype, source, tag, communicator, &status);`
- Fortran
  - `Call MPI_RECV(buffer, count, datatype, source, tag, communicator, status, ierr)`
- Call blocks the program until message is in buffer
- Status - contains information about incoming message
  - C
    - `MPI_Status status;`
  - Fortran
    - `Integer status(MPI_STATUS_SIZE)`



**Call MPI\_Recv(buffer, count, datatype,  
source, tag, communicator,  
status, ierr)**

- **Buffer**: The data array to be received
- **Count** : Maximum length of data array  
(in elements, 1 for scalars)
- **Datatype** : Type of data, for example :  
MPI\_DOUBLE\_PRECISION, MPI\_INT, etc
- **Source** : Source processor number  
(within given communicator)
- **Tag** : Message type (arbitrary integer)
- **Communicator** : Your set of processors
- **Status**: Information about message
- **Ierr** : Error return (Fortran only)

# Exercise 2 : Basic Send and Receive

- Write a parallel program to send & receive data
  - Initialize MPI
  - Have processor 0 send an integer to processor 1
  - Have processor 1 receive an integer from processor 0
  - Both processors print the data
  - Quit MPI



# Send and Recv in MPI

```
el2:mpi> cat f_ex01.f90
      module fmpi
!DEC$ NOFREEFORM
      include "mpif.h"
!DEC$ FREEFORM
      end module

!*****
! This is a simple send/receive program in MPI
! Processor 0 sends an integer to processor 1,
! while processor 1 receives the integer from proc. 0
!*****
      program hello
      use fmpi
!      include "mpif.h"
      integer myid, ierr,numprocs
      integer tag,source,destination,count
      integer buffer
      integer status(MPI_STATUS_SIZE)
      call MPI_INIT( ierr )
      call MPI_COMM_RANK( MPI_COMM_WORLD, myid, ierr )
      call MPI_COMM_SIZE( MPI_COMM_WORLD, numprocs, ierr )
      tag=1234
      source=0
      destination=1
      count=1
      if(myid .eq. source)then
        buffer=5678
        Call MPI_Send(buffer, count, MPI_INTEGER,destination,&
          tag, MPI_COMM_WORLD, ierr)
        write(*,*)"processor ",myid," sent ",buffer
      endif
      if(myid .eq. destination)then
        Call MPI_Recv(buffer, count, MPI_INTEGER,source,&
          tag, MPI_COMM_WORLD, status,ierr)
        write(*,*)"processor ",myid," got ",buffer
      endif
      call MPI_FINALIZE(ierr)
      stop
      end
```

# Send an Recv in MPI

```
el2:mpi> cat c_ex01.c
#include <stdio.h>
#include <stdlib.h>
#include <mpi.h>
#include <math.h>
```

```

/*****
This is a simple send/receive program in MPI
*****/
```

```
int main(argc,argv)
int argc;
char *argv[];
{
    int myid, numprocs;
    int tag,source,destination,count;
    int buffer;
    MPI_Status status;

    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);
    tag=1234;
    source=0;
    destination=1;
    count=1;
    if(myid == source){
        buffer=5678;
        MPI_Send(&buffer,count,MPI_INT,destination,tag,MPI_COMM_WORLD);
        printf("processor %d  sent %d\n",myid,buffer);
    }
    if(myid == destination){
        MPI_Recv(&buffer,count,MPI_INT,source,tag,MPI_COMM_WORLD,&status);
        printf("processor %d  got %d\n",myid,buffer);
    }
    MPI_Finalize();
}
```

```
el2:mpi> cat slurm-7068778.out
processor          0  sent          5678
processor          1  got           5678
el2:mpi>
```



# Summary

- MPI is used to create parallel programs based on message passing
- Usually the same program is run on multiple processors
- Well over 100 "Advanced Calls"
- The 6 basic calls in MPI are:

– `MPI_INIT( ierr )`

– `MPI_COMM_RANK( MPI_COMM_WORLD, myid, ierr )`

– `MPI_COMM_SIZE( MPI_COMM_WORLD, numprocs, ierr )`

– `MPI_Send(buffer, count, MPI_INTEGER, destination, tag, MPI_COMM_WORLD, ierr)`

– `MPI_Recv(buffer, count, MPI_INTEGER, source, tag, MPI_COMM_WORLD, status, ierr)`

– `MPI_FINALIZE(ierr)`

# MPI is not limited to huge HPC

- Can be built on laptops and desktops
  - Might be useful for testing
  - Might actually get good performance
- I have it on my laptop & high end desktop
- I have it running on my home 4 node Raspberry Pi cluster
  - <https://www.raspberrypi.org>
  - < \$100/node



# Unofficial Languages

- Subsets available for R and Java
- Fairly complete implementation for Python
- Let me know if you are interested in these
- I have not looked at Julia

# Java

- Comes in two versions OpenMPI and Intel
- They are not compatible even at the source (API) level
- `cd ../java`
- `diff Hello.java lhello.java`

```
...  
<  int myrank = MPI.COMM_WORLD.getRank();  
<  int size = MPI.COMM_WORLD.getSize() ;  
>  mpi.Comm  comm = mpi.Comm.WORLD;  
>  int myrank = comm.getRank();  
>  int size = comm.getSize();  
  
<  MPI.COMM_WORLD.send(message, tosend, MPI.INT, 1, tag);  
>  mpi.PTP.send(message,10,mpi.Datatype.INT,1,tag,comm);  
  
<  MPI.COMM_WORLD.recv(message, toget, MPI.INT, 0, tag);  
>  mpi.PTP.recv(message,10,mpi.Datatype.INT,0,tag,comm);
```



# Build Rmpi & MPI4py

I used a Conda environment  
that already had R and  
Python. Then I compiled  
using IntelMPI, not OpenMPI

```
cd mixedlang
```

```
ml conda  
conda activate plex1  
which python  
which R  
ml comp-intel/2020.1.217  
ml intel-mpi/2020.1.217  
which mpicc
```

```
curl --insecure https://cran.r-project.org/src/contrib/Rmpi_0.6-9.2.tar.gz -o Rmpi.tar.gz
```

```
export TYPE=OPENMPI  
export MY_MPI_PATH=/nopt/nrel/apps/compilers/intel/2020.1.217/impi/2019.7.217/intel64
```

```
R CMD INSTALL --configure-args="\n--with-Rmpi-include='$MY_MPI_PATH/include' \n--with-Rmpi-libpath='$MY_MPI_PATH/lib/release' \n--with-mpi='$MY_MPI_PATH/bin/mpicc' \n--with-Rmpi-type='$TYPE'" Rmpi.tar.gz
```

```
pip install mpi4py
```

# It is possible(sometimes) to mix languages

```
(/home/tkaiser2/.conda-envs/plex1) el3:mixedlang> cat together
#!/bin/bash
#SBATCH --job-name="flow"
#SBATCH --nodes=1
#SBATCH --export=ALL
#SBATCH --oversubscribe
#SBATCH --time=00:10:00
#SBATCH --partition=debug
#SBATCH --account=hpcapps
#
##r_ex01c.R does not work with openmpi and other programs
#j_ex01c.java only works with openmpi
#ij_ex01c.java is for Intel MPI and works with all others
#
#
export PATH=/nopt/nrel/apps/openmpi/4.1.0-gcc-8.4.0-j15/jdk-15.0.2/bin:$PATH
module purge
ml conda
conda activate plex1
which python
which R
ml comp-intel/2020.1.217
ml intel-mpi/2020.1.217
which mpicc

mpif90 f_ex01c.f90 -o f_ex01c ; rm -rf fmpi.mod
mpicc c_ex01c.c -o c_ex01c

#for source in f_ex01c c_ex01c r_ex01c.R P_ex01c.py "java ij_ex01c.java" ; do
#for source in f_ex01c c_ex01c r_ex01c.R P_ex01c.py ; do
for source in f_ex01c c_ex01c r_ex01c.R P_ex01c.py ; do
    echo DRIVER: $source
#mpiexec -n 1 $source 2 5 : -n 1 ./f_ex01c : -n 1 ./c_ex01c : -n 1 ./r_ex01c.R : -n 1 ./P_ex01c.py : -n 1 java ij_ex01c.java
| egrep "got"
    echo 0 ./ $source > mapfile
    echo 1 ./f_ex01c >> mapfile
    echo 2 ./c_ex01c >> mapfile
    echo 3 ./r_ex01c.R >> mapfile
    echo 4 ./P_ex01c.py >> mapfile
    echo 5 java ./ij_ex01c.java >> mapfile
#srun -n 5 --multi-prog mapfile | grep got
srun --partition=debug --time=00:10:00 -n 6 --multi-prog mapfile | grep got
    echo " "
done
```



DRIVER: f\_ex01c

Fortran processor	1	got	4678	3678	2678	1678
Fortran processor	1	got	4679	3679	2679	1679
Fortran processor	1	got	4680	3680	2680	1680
C processor 2	got	4678	3678	2678	1678	
C processor 2	got	4679	3679	2679	1679	
C processor 2	got	4680	3680	2680	1680	
Java processor 5	got	4678	3678	2678	1678	
Java processor 5	got	4679	3679	2679	1679	
Java processor 5	got	4680	3680	2680	1680	
[1] "R processor	3	got	4678	3678	2678	1678"
[1] "R processor	3	got	4679	3679	2679	1679"
[1] "R processor	3	got	4680	3680	2680	1680"
Python processor	4	got	[4678	3678	2678	1678]
Python processor	4	got	[4679	3679	2679	1679]
Python processor	4	got	[4680	3680	2680	1680]

DRIVER: c\_ex01c

C processor 2	got	1	101	201	301		
Fortran processor	1	got		1	101	201	301
Fortran processor	1	got		2	102	202	302
Fortran processor	1	got		3	103	203	303
C processor 2	got	2	102	202	302		
C processor 2	got	3	103	203	303		
Java processor 5	got	1	101	201	301		
Java processor 5	got	2	102	202	302		
Java processor 5	got	3	103	203	303		
[1] "R processor	3	got	1	101	201	301"	
[1] "R processor	3	got	2	102	202	302"	
[1] "R processor	3	got	3	103	203	303"	
Python processor	4	got	[	1	101	201	301]
Python processor	4	got	[	2	102	202	302]
Python processor	4	got	[	3	103	203	303]

DRIVER: r\_ex01c.R

Fortran processor	1	got		3	5	7	9
C processor 2	got	3	5	7	9		
Fortran processor	1	got		4	6	8	10
C processor 2	got	4	6	8	10		
Fortran processor	1	got		5	7	9	11
C processor 2	got	5	7	9	11		
[1] "R processor	3	got	3	5	7	9"	
[1] "R processor	3	got	4	6	8	10"	
[1] "R processor	3	got	5	7	9	11"	
Java processor 5	got	3	5	7	9		
Java processor 5	got	4	6	8	10		
Java processor 5	got	5	7	9	11		
Python processor	4	got	[3	5	7	9]	
Python processor	4	got	[	4	6	8	10]
Python processor	4	got	[	5	7	9	11]

DRIVER: P\_ex01c.py

Fortran processor	1	got	6679	7679	8679	9679
C processor 2	got	6679	7679	8679	9679	
Fortran processor	1	got	6680	7680	8680	9680
C processor 2	got	6680	7680	8680	9680	
Fortran processor	1	got	6681	7681	8681	9681
C processor 2	got	6681	7681	8681	9681	
Java processor 5	got	6679	7679	8679	9679	
Java processor 5	got	6680	7680	8680	9680	
Java processor 5	got	6681	7681	8681	9681	
[1] "R processor	3	got	6679	7679	8679	9679"
[1] "R processor	3	got	6680	7680	8680	9680"
[1] "R processor	3	got	6681	7681	8681	9681"
Python processor	4	got	[6679	7679	8679	9679]
Python processor	4	got	[6680	7680	8680	9680]
Python processor	4	got	[6681	7681	8681	9681]