

Parallel Programming

Basic MPI

Timothy H. Kaiser, Ph.D.
tkaiser2@nrel.gov

Session 1: Introduction to MPI (this invitation)

Monday, January 24th 11:00 AM - 12:00 PM Mountain

The first session will introduce MPI. We will give a background, show some sources of Documentation. We will show the classic "Hello world" program in MPI running on multiple processors. We will discuss Basic communications and show a simple send and receive program where messages are passed between processors. We will be running examples in the official languages supported by MPI, C and Fortran. Also, we will briefly discuss support for Python, R, and Java. Source code and scripts will be provided that can be run on Eagle.

Session 2: Expansion to Higher-level MPI Calls

Monday, January 31st 11:00 AM - 12:00 PM Mountain

The second session will expand on the first, showing many of the higher-level MPI calls commonly used to write parallel programs. Examples will be provided that can run on Eagle. We will look at: using the various predefined data types, broadcast, wildcards, asynchronous communications, and using probes and status information to control flow.

Session 3: Additional Collective Operations

Monday, February 7th 11:00 AM - 12:00 PM Mountain

In the third MPI session, we will look at additional collective operations including scatter, gather, and reductions. We'll also show examples of the "variable" versions of these calls where the amount of information shared is processor-dependent. We'll look at creating derived data types and managing subsets of processes using communicators.

Session 4: Finite Difference Model

Monday, February 14th 11:00 AM - 12:00 PM Mountain

In the fourth session, we will introduce a finite difference model that will demonstrate what a computational scientist needs to do to take advantage of computers using MPI. The model we are using is a two-dimensional finite-difference code. After discussing the serial code, we will show the modifications necessary to turn it into a parallel program using MPI. We will look at domain decomposition, initialization, data distribution, message passing, reduction operations, and multiple methods for data output. We will also look at the performance of the application on various numbers of processors to illustrate Amdahl's parallel program scaling law.

Examples and Slides

Examples:

```
git clone https://github.com/timkphd/examples  
cd examples/mpi
```

Slides:

```
https://github.com/timkphd/slides
```


Talk Overview

- Background on MPI
- Documentation
- Hello world in MPI
- Basic communications
- Simple send and receive program

Background on MPI

- MPI - Message Passing Interface
 - Library standard defined by a committee of vendors, implementers, & parallel programmers
 - Used to create **parallel programs** based on message passing
- Portable: one standard, many implementations
- Available on almost all parallel machines in **C and Fortran**
- Over 100 advanced routines but 6 basic
- Also used to create **libraries**.

Unofficial Languages

- Subsets available for R and Java
- Fairly complete implementation for Python
- Let me know if you are interested in these
- I have not looked at Julia

Documentation

- MPI home page (contains the library standard)
<https://www.mpi-forum.org>
- Books
 - "MPI: The Complete Reference" by Snir, Otto, Huss-Lederman, Walker, and Dongarra, MIT Press (also in Postscript and html)
 - "Using MPI" by Gropp, Lusk and Skjellum, MIT Press
- Tutorials
 - many online, just do a search

MPI Implementations

- Most parallel supercomputer vendors provide optimized implementations
- Intel
- HP/SGI/Cray - mpt
- IBM

Open Source MPI Implementations

- MPICH:
 - <https://www.mpich.org>
- Mvapich
 - <http://mvapich.cse.ohio-state.edu>
- OpenMPI
 - www.open-mpi.org

Key Concepts of MPI

- Used to create parallel programs based on message passing
- Normally the same program is running on several different processors
- Processors **communicate** using **message** passing
- Typical methodology:

```
start job on n processors
do i=1 to j
  each processor does some calculation
  pass messages between processor
end do
end job
```

Communicators

- Communicator
 - A collection of processors working on some part of a parallel job
 - Used as a parameter for most MPI calls
 - `MPI_COMM_WORLD` includes all of the processors in your job
 - Processors within a communicator are assigned numbers (ranks) 0 to $n-1$
 - Can create subsets of `MPI_COMM_WORLD`

Messages

- Simplest message: an array of data of one type.
- Predefined types correspond to commonly used types in a given language
 - MPI_REAL (Fortran), MPI_FLOAT (C)
 - MPI_DOUBLE_PRECISION (Fortran), MPI_DOUBLE (C)
 - MPI_INTEGER (Fortran), MPI_INT (C)
- User can define more complex types and send packages.

Include files

- The MPI include file
 - C: `mpi.h`
 - Fortran: `mpif.h` (a f90 module is a good place for this)
- Defines many constants used within MPI programs
- In C defines the interfaces for the functions
- Compilers know where to find the include files

Minimal MPI program

- Every MPI program needs these...
- C version

```
/* the mpi include file */
#include <mpi.h>
int nPES, ierr, iam;
/* Initialize MPI */
ierr=MPI_Init(&argc, &argv);
/* How many processors (nPES) are there? */
ierr=MPI_Comm_size(MPI_COMM_WORLD, &nPES);
/* What processor am I (what is my rank)? */
ierr=MPI_Comm_rank(MPI_COMM_WORLD, &iam);
...
ierr=MPI_Finalize();
```

In C MPI routines are functions and return an error value

Minimal MPI program

- Every MPI program needs these...
- Fortran version

```
! MPI include file
  include 'mpif.h'
! The mpi module can be used for Fortran 90 instead of mpif.h
!   use mpi
  integer nPEs, ierr, iam
! Initialize MPI
  call MPI_Init(ierr)
! How many processors (nPEs) are there?
  call MPI_Comm_size(MPI_COMM_WORLD, nPEs, ierr)
! What processor am I (what is my rank)?
  call MPI_Comm_rank(MPI_COMM_WORLD, iam, ierr)
  ...
  call MPI_Finalize(ierr)
```

In Fortran, MPI routines are subroutines, and last parameter is an error value

Exercise I : Hello World

- Write a parallel “hello world” program
 - Initialize MPI
 - Have each processor print out “Hello, World” and its processor number (rank)
 - Quit MPI

Fortran and C examples

```
el2:mpi> cat helloc.c
#include <stdio.h>
#include <stdlib.h>
#include <mpi.h>
#include <math.h>

/*****
This is a simple hello world program. Each processor prints
name, rank, and total run size.
*****/
int main(int argc, char **argv)
{
    int myid,numprocs,resultlen;
    char myname[MPI_MAX_PROCESSOR_NAME] ;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);

    MPI_Get_processor_name(myname,&resultlen);

    printf("Hello from %s %d %d\n",myname,myid,numprocs);
    MPI_Finalize();
}
```


Fortran and C examples

```
el2:mpi> cat hellof.f90
!*****
! This is a simple hello world program. Each processor
! prints out its name, rank and number of processors
! in the current MPI run.
!*****
  program hello
  include "mpif.h"
  integer myid,numprocs,ierr,nlength
  character (len=MPI_MAX_PROCESSOR_NAME):: myname
  call MPI_INIT( ierr )
  call MPI_COMM_RANK( MPI_COMM_WORLD, myid, ierr )
  call MPI_COMM_SIZE( MPI_COMM_WORLD, numprocs, ierr )

  call MPI_Get_processor_name(myname,nlength,ierr)

  write (*,*) "Hello from ",trim(myname)," # ",myid," of ",numprocs
  call MPI_FINALIZE(ierr)
  stop
end
el2:mpi>
el2:mpi>
```

Compiling

- Most MPI compilers are actually just scripts that call underlying Fortran or C compilers
- Load module that points to your compilers
- For Intel compilers use
 - `module load intel-mpi`
 - `mpiifort`, `mpiicc`, `mpicpc` - Intel backend compilers
 - `mpif90` `mpicc`, `mpicxx` - gcc/gfortran backend compilers
- For HP's version of MPI
 - `module load mpt gcc/8.4.0`
 - `mpif77` `mpif90`
 - `mpicc` `mpiCC`

Compiling IntelMPI and MPT

```
el2:mpi> ml intel-mpi
el2:mpi>
el2:mpi> mpiicc helloc.c -o helloc
el2:mpi> mpiifort hellof.f90 -o hellof
el2:mpi> ##### or #####
el2:mpi> mpicc helloc.c -o helloc
el2:mpi> mpif90 hellof.f90 -o hellof
```

Or

```
el2:mpi> module load mpt gcc/8.4.0
el2:mpi>
el2:mpi> mpicc helloc.c -o hello
el2:mpi> mpif90 hellof.f90 -o hellof
el2:mpi>
```

Compiling OpenMPI

- Most MPI compilers are actually just scripts that call underlying Fortran or C compilers
- Load module that points to your compilers
- Openmpi
 - `module load openmpi`
 - `mpif77 mpif90`
 - `mpicc mpiCC mpic++ mpicxx`
- This will give you gcc/gfortran as the backend compilers
- To get Intel backend compilers
 - `export OMPI_FC=ifort`
 - `export OMPI_CC=icc`
 - `export OMPI_CXX=icpc`

makefile

```
el1:mpi> ml intel-mpi
```

```
el1:mpi> make
```

```
mpicc      c_ex00.c -o c_ex00
```

```
mpicc      c_ex01.c -o c_ex01
```

```
mpicc      c_ex02.c -o c_ex02
```

```
mpicc      c_ex03.c -o c_ex03
```

```
mpicc      c_ex04.c -o c_ex04
```

```
mpicc      c_ex05.c -o c_ex05
```

```
mpicc      c_ex06.c -o c_ex06
```

```
mpicc      c_ex07.c -o c_ex07
```

```
mpicc      c_ex08.c -o c_ex08
```

```
mpicc      c_ex09.c -o c_ex09
```

```
mpicc      c_ex10.c -o c_ex10
```

```
mpicc      c_ex11.c -o c_ex11
```

```
mpicc      c_ex12.c -o c_ex12
```

```
mpicc      c_ex13.c -o c_ex13
```

```
c_ex13.c: In function 'main':
```

```
c_ex13.c:34:23: warning: cast to pointer from integer of different size [-Wint-to-pointer-cast]
```

```
    mpi_err = MPI_Gather((void*)mysize,1,MPI_INT,
```

```
mpicc      helloc.c -o helloc
```

```
mpif90     f_ex00.f90 -o f_ex00
```

```
mpif90     f_ex01.f90 -o f_ex01
```

```
mpif90     f_ex02.f90 -o f_ex02
```

```
mpif90     f_ex03.f90 -o f_ex03
```

```
mpif90     f_ex04.f90 -o f_ex04
```

```
mpif90     f_ex05.f90 -o f_ex05
```

```
mpif90     f_ex06.f90 -o f_ex06
```

```
mpif90     f_ex07.f90 -o f_ex07
```

```
mpif90     f_ex08.f90 -o f_ex08
```

```
mpif90     f_ex09.f90 -o f_ex09
```

```
mpif90     f_ex10.f90 -o f_ex10
```

```
mpif90     f_ex11.f90 -o f_ex11
```

```
mpif90     f_ex12.f90 -o f_ex12
```

```
mpif90     f_ex13.f90 -o f_ex13
```

```
mpif90     hellof.f90 -o hellof
```

```
el1:mpi>
```

Running

- Most often you will use a batch system
- Write a batch script file.
- Use the command ~~**mpiexec**~~ or ~~**mpirun**~~ **srun** to actually start the program. You must tell the system how many copies to run
- On some systems you must tell where to run the program

A very simple Slurm Script - "sim"

```
#!/bin/bash
#SBATCH --job-name="flow"
#SBATCH --nodes=1
#SBATCH --export=ALL
#SBATCH --oversubscribe
#SBATCH --time=00:10:00
#SBATCH --partition=debug
#SBATCH --account=hpcapps
```

```
module purge
ml intel-mpi
```

```
srun -n 4 ./helloc
```

The run...

```
el1:mpi> sbatch --account=hpcapps sim
Submitted batch job 8182395
el1:mpi> cat slurm-8182395.out
Hello from r1i7n35 #      2 of      4
Hello from r1i7n35 #      0 of      4
Hello from r1i7n35 #      1 of      4
Hello from r1i7n35 #      3 of      4
el1:mpi>
```

A useful example - "doeagle"


```
#!/bin/bash
#SBATCH --job-name="flow"
#SBATCH --nodes=1
#SBATCH --export=ALL
#SBATCH --oversubscribe
#SBATCH --time=00:10:00
#SBATCH --partition=debug
#SBATCH --account=hpcapps
```

```
module purge
ml intel-mpi
```

```
if [ -z ${EXE+x} ]; then export EXE=c_ex00 ; fi
```

```
echo "running " $EXE
srun -n 4 ./ $EXE
```

You can set the
program to run with
and export before
you submit the run



```
el1:mpi> export EXE=f_ex02
el1:mpi> sbatch --account=hpcapps doeagle
Submitted batch job 8182393
el1:mpi> cat *8182393*
running f_ex02
Hello from fortran process:
Hello from fortran process:
Hello from fortran process:
Hello from fortran process:
getting 1
i= 100
el1:mpi>
```

The run...

3	Numprocs is	4
1	Numprocs is	4
2	Numprocs is	4
0	Numprocs is	4

Slight Digression

- One way you can understand what is going on in a MPI application is to look at it running in a debugger
- We have ARM/ddt
- We also have ARM/map and ARM/Profile for profiling
- Let me know if interested

ARM ddt

Arm DDT - Arm Forge 21.1.2

Current Group: All Focus on current: ☒ Group ☐ Process ☐ Thread ☐ Step Threads Together

All 0 1 2 3

Create Group

Project Files

Search (%K)

Application Code

Sources

c_ex00.c

main(arg)

External Code

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <mpi.h>
4 #include <math.h>
5
6 /*****
7 This is a simple hello world program. Each processor prints out
8 it's rank and the size of the current MPI run (Total number of
9 processors).
10 *****/
11 int main(argc,argv)
12 int argc;
13 char *argv[];
14 {
15     int myid, numprocs;
16
17     MPI_Init(&argc,&argv);
18     MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
19     MPI_Comm_rank(MPI_COMM_WORLD,&myid);
20
21
22     /* print out my rank and this run's PE size */
23     printf("Hello from task %d of %d\n",myid,numprocs);
24
25     MPI_Finalize();
26 }
27
28
29
```

Locals Current Line(s) Current Stack

Current Line(s)

Name	Value
numprocs	0

Breakp... Watchp... Stacks Tracep... Tracepoint O...

Stacks

Processes	Threads	Function
4	4	main (c_ex00.c:18)
4	4	ucs_async_thread_func

Evaluate

Name	Value
------	-------

Ready Connected to: (via tunnel) el1:4201 -> el1

Basic Communication

- Data values are transferred from one processor to another
 - One processor sends the data
 - Another receives the data
- Synchronous
 - Call does not return until the message is sent or received
- Asynchronous
 - Call indicates a start of send or receive, and another call is made to determine if finished

Synchronous Send

- C
 - `MPI_Send(&buffer, count, datatype, destination, tag, communicator);`
- Fortran
 - `Call MPI_Send(buffer, count, datatype, destination, tag, communicator, ierr)`
 - Call blocks until message on the way

**Call MPI_Send(buffer, count, datatype,
destination, tag, communicator, ierr)**

- **Buffer**: The data array to be sent
- **Count** : Length of data array (in elements, 1 for scalars)
- **Datatype** : Type of data, for example : MPI_DOUBLE_PRECISION, MPI_INT, etc
- **Destination** : Destination processor number (within given communicator)
- **Tag** : Message type (arbitrary integer)
- **Communicator** : Your set of processors
- **Ierr** : Error return (Fortran only)

Synchronous Receive

- C
 - `MPI_Recv(&buffer, count, datatype, source, tag, communicator, &status);`
- Fortran
 - `Call MPI_RECV(buffer, count, datatype, source, tag, communicator, status, ierr)`
- Call blocks the program until message is in buffer
- Status - contains information about incoming message
 - C
 - `MPI_Status status;`
 - Fortran
 - `Integer status(MPI_STATUS_SIZE)`

**Call MPI_Recv(buffer, count, datatype,
source, tag, communicator,
status, ierr)**

- **Buffer**: The data array to be received
- **Count** : Maximum length of data array
(in elements, 1 for scalars)
- **Datatype** : Type of data, for example :
MPI_DOUBLE_PRECISION, MPI_INT, etc
- **Source** : Source processor number
(within given communicator)
- **Tag** : Message type (arbitrary integer)
- **Communicator** : Your set of processors
- **Status**: Information about message
- **Ierr** : Error return (Fortran only)

Exercise 2 : Basic Send and Receive

- Write a parallel program to send & receive data
 - Initialize MPI
 - Have processor 0 send an integer to processor 1
 - Have processor 1 receive an integer from processor 0
 - Both processors print the data
 - Quit MPI

Send and Recv in MPI

```
el2:mpi> cat f_ex01.f90
      module fmpi
!DEC$ NOFREEFORM
      include "mpif.h"
!DEC$ FREEFORM
      end module

!*****
!  This is a simple send/receive program in MPI
!  Processor 0 sends an integer to processor 1,
!  while processor 1 receives the integer from proc. 0
!*****

      program hello
      use fmpi
!      include "mpif.h"
      integer myid, ierr,numprocs
      integer tag,source,destination,count
      integer buffer
      integer status(MPI_STATUS_SIZE)
      call MPI_INIT( ierr )
      call MPI_COMM_RANK( MPI_COMM_WORLD, myid, ierr )
      call MPI_COMM_SIZE( MPI_COMM_WORLD, numprocs, ierr )
      tag=1234
      source=0
      destination=1
      count=1
      if(myid .eq. source)then
        buffer=5678
        Call MPI_Send(buffer, count, MPI_INTEGER,destination,&
          tag, MPI_COMM_WORLD, ierr)
        write(*,*)"processor ",myid," sent ",buffer
      endif
      if(myid .eq. destination)then
        Call MPI_Recv(buffer, count, MPI_INTEGER,source,&
          tag, MPI_COMM_WORLD, status,ierr)
        write(*,*)"processor ",myid," got ",buffer
      endif
      call MPI_FINALIZE(ierr)
      stop
      end
```

Send an Recv in MPI

```
el2:mpi> cat c_ex01.c
#include <stdio.h>
#include <stdlib.h>
#include <mpi.h>
#include <math.h>
```

```
/******
This is a simple send/receive program in MPI
*****/
```

```
int main(argc,argv)
int argc;
char *argv[];
{
    int myid, numprocs;
    int tag,source,destination,count;
    int buffer;
    MPI_Status status;

    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);
    tag=1234;
    source=0;
    destination=1;
    count=1;
    if(myid == source){
        buffer=5678;
        MPI_Send(&buffer,count,MPI_INT,destination,tag,MPI_COMM_WORLD);
        printf("processor %d  sent %d\n",myid,buffer);
    }
    if(myid == destination){
        MPI_Recv(&buffer,count,MPI_INT,source,tag,MPI_COMM_WORLD,&status);
        printf("processor %d  got %d\n",myid,buffer);
    }
    MPI_Finalize();
}
```

```
el2:mpi> cat slurm-7068778.out
processor          0  sent          5678
processor          1  got           5678
el2:mpi>
```


Summary

- MPI is used to create parallel programs based on message passing
- Usually the same program is run on multiple processors
- Well over 100 "Advanced Calls"
- The 6 basic calls in MPI are:

– `MPI_INIT(ierr)`

– `MPI_COMM_RANK(MPI_COMM_WORLD, myid, ierr)`

– `MPI_COMM_SIZE(MPI_COMM_WORLD, numprocs, ierr)`

– `MPI_Send(buffer, count, MPI_INTEGER, destination, tag, MPI_COMM_WORLD, ierr)`

– `MPI_Recv(buffer, count, MPI_INTEGER, source, tag, MPI_COMM_WORLD, status, ierr)`

– `MPI_FINALIZE(ierr)`

MPI is not limited to huge HPC

- Can be built on laptops and desktops
 - Might be useful for testing
 - Might actually get good performance
- I have it on my laptop & high end desktop
- I have it running on my home 4 node Raspberry Pi cluster
 - <https://www.raspberrypi.org>
 - < \$100/node

Unofficial Languages

- Subsets available for R and Java
- Fairly complete implementation for Python
- Let me know if you are interested in these
- I have not looked at Julia

Java

- Comes in two versions OpenMPI and Intel
- They are not compatible even at the source (API) level
- `cd ../java`
- `diff Hello.java lhello.java`

```
...  
<  int myrank = MPI.COMM_WORLD.getRank();  
<  int size = MPI.COMM_WORLD.getSize() ;  
>  mpi.Comm comm = mpi.Comm.WORLD;  
>  int myrank = comm.getRank();  
>  int size = comm.getSize();  
  
<  MPI.COMM_WORLD.send(message, tosend, MPI.INT, 1, tag);  
>  mpi.PTP.send(message,10,mpi.Datatype.INT,1,tag,comm);  
  
<  MPI.COMM_WORLD.recv(message, toget, MPI.INT, 0, tag);  
>  mpi.PTP.recv(message,10,mpi.Datatype.INT,0,tag,comm);
```


Build Rmpi & MPI4py

I used a Conda environment
that already had R and
Python. Then I compiled
using IntelMPI, not OpenMPI

```
cd mixedlang
```

```
ml conda  
conda activate plex1  
which python  
which R  
ml comp-intel/2020.1.217  
ml intel-mpi/2020.1.217  
which mpicc
```

```
curl --insecure https://cran.r-project.org/src/contrib/Rmpi_0.6-9.2.tar.gz -o Rmpi.tar.gz
```

```
export TYPE=OPENMPI  
export MY_MPI_PATH=/nopt/nrel/apps/compilers/intel/2020.1.217/impi/2019.7.217/intel64
```

```
R CMD INSTALL --configure-args="\n--with-Rmpi-include='$MY_MPI_PATH/include' \n--with-Rmpi-libpath='$MY_MPI_PATH/lib/release' \n--with-mpi='$MY_MPI_PATH/bin/mpicc' \n--with-Rmpi-type='$TYPE'" Rmpi.tar.gz
```

```
pip install mpi4py
```

It is possible(sometimes) to mix languages

```
(/home/tkaiser2/.conda-envs/plex1) el3:mixedlang> cat together
#!/bin/bash
#SBATCH --job-name="flow"
#SBATCH --nodes=1
#SBATCH --export=ALL
#SBATCH --oversubscribe
#SBATCH --time=00:10:00
#SBATCH --partition=debug
#SBATCH --account=hpcapps
#
##r_ex01c.R does not work with openmpi and other programs
#j_ex01c.java only works with openmpi
#ij_ex01c.java is for Intel MPI and works with all others
#
#
export PATH=/nopt/nrel/apps/openmpi/4.1.0-gcc-8.4.0-j15/jdk-15.0.2/bin:$PATH
module purge
ml conda
conda activate plex1
which python
which R
ml comp-intel/2020.1.217
ml intel-mpi/2020.1.217
which mpicc

mpif90 f_ex01c.f90 -o f_ex01c ; rm -rf fmpi.mod
mpicc c_ex01c.c -o c_ex01c

#for source in f_ex01c c_ex01c r_ex01c.R P_ex01c.py "java ij_ex01c.java" ; do
#for source in f_ex01c c_ex01c r_ex01c.R P_ex01c.py ; do
for source in f_ex01c c_ex01c r_ex01c.R P_ex01c.py ; do
    echo DRIVER: $source
#mpiexec -n 1 $source 2 5 : -n 1 ./f_ex01c : -n 1 ./c_ex01c : -n 1 ./r_ex01c.R : -n 1 ./P_ex01c.py : -n 1 java ij_ex01c.java
| egrep "got"
    echo 0 ./ $source > mapfile
    echo 1 ./f_ex01c >> mapfile
    echo 2 ./c_ex01c >> mapfile
    echo 3 ./r_ex01c.R >> mapfile
    echo 4 ./P_ex01c.py >> mapfile
    echo 5 java ./ij_ex01c.java >> mapfile
#srun -n 5 --multi-prog mapfile | grep got
srun --partition=debug --time=00:10:00 -n 6 --multi-prog mapfile | grep got
    echo " "
done
```


DRIVER: f_ex01c

Fortran processor	1	got	4678	3678	2678	1678
Fortran processor	1	got	4679	3679	2679	1679
Fortran processor	1	got	4680	3680	2680	1680
C processor 2	got	4678	3678	2678	1678	
C processor 2	got	4679	3679	2679	1679	
C processor 2	got	4680	3680	2680	1680	
Java processor 5	got	4678	3678	2678	1678	
Java processor 5	got	4679	3679	2679	1679	
Java processor 5	got	4680	3680	2680	1680	
[1] "R processor	3	got	4678	3678	2678	1678"
[1] "R processor	3	got	4679	3679	2679	1679"
[1] "R processor	3	got	4680	3680	2680	1680"
Python processor	4	got	[4678	3678	2678	1678]
Python processor	4	got	[4679	3679	2679	1679]
Python processor	4	got	[4680	3680	2680	1680]

DRIVER: c_ex01c

C processor 2	got	1	101	201	301		
Fortran processor	1	got		1	101	201	301
Fortran processor	1	got		2	102	202	302
Fortran processor	1	got		3	103	203	303
C processor 2	got	2	102	202	302		
C processor 2	got	3	103	203	303		
Java processor 5	got	1	101	201	301		
Java processor 5	got	2	102	202	302		
Java processor 5	got	3	103	203	303		
[1] "R processor	3	got	1	101	201	301"	
[1] "R processor	3	got	2	102	202	302"	
[1] "R processor	3	got	3	103	203	303"	
Python processor	4	got	[1	101	201	301]
Python processor	4	got	[2	102	202	302]
Python processor	4	got	[3	103	203	303]

DRIVER: r_ex01c.R

Fortran processor	1	got		3	5	7	9
C processor 2	got	3	5	7	9		
Fortran processor	1	got		4	6	8	10
C processor 2	got	4	6	8	10		
Fortran processor	1	got		5	7	9	11
C processor 2	got	5	7	9	11		
[1] "R processor	3	got	3	5	7	9"	
[1] "R processor	3	got	4	6	8	10"	
[1] "R processor	3	got	5	7	9	11"	
Java processor 5	got	3	5	7	9		
Java processor 5	got	4	6	8	10		
Java processor 5	got	5	7	9	11		
Python processor	4	got	[3	5	7	9]	
Python processor	4	got	[4	6	8	10]
Python processor	4	got	[5	7	9	11]

DRIVER: P_ex01c.py

Fortran processor	1	got	6679	7679	8679	9679
C processor 2	got	6679	7679	8679	9679	
Fortran processor	1	got	6680	7680	8680	9680
C processor 2	got	6680	7680	8680	9680	
Fortran processor	1	got	6681	7681	8681	9681
C processor 2	got	6681	7681	8681	9681	
Java processor 5	got	6679	7679	8679	9679	
Java processor 5	got	6680	7680	8680	9680	
Java processor 5	got	6681	7681	8681	9681	
[1] "R processor	3	got	6679	7679	8679	9679"
[1] "R processor	3	got	6680	7680	8680	9680"
[1] "R processor	3	got	6681	7681	8681	9681"
Python processor	4	got	[6679	7679	8679	9679]
Python processor	4	got	[6680	7680	8680	9680]
Python processor	4	got	[6681	7681	8681	9681]