



A Prototype Finite Difference Model

Timothy H. Kaiser, Ph.D.



Slides:

<https://github.com/timkphd/slides>

Source

```
hexagon:~ tkaiser$ mkdir tut
hexagon:~ tkaiser$ cd tut
hexagon:tut tkaiser$ git clone https://github.com/timkphd/examples.git
Cloning into 'examples'...
remote: Enumerating objects: 83, done.
remote: Counting objects: 100% (83/83), done.
remote: Compressing objects: 100% (63/63), done.
remote: Total 1236 (delta 39), reused 45 (delta 19), pack-reused 1153
Receiving objects: 100% (1236/1236), 2.47 MiB | 3.64 MiB/s, done.
Resolving deltas: 100% (390/390), done.
hexagon:tut tkaiser$ mv examples/r .
hexagon:tut tkaiser$ rm -rf examples
hexagon:tut tkaiser$ cd r

hexagon:r tkaiser$ ./getdat.py
downloading file
True
split -l 158563 start start
hexagon:r tkaiser$ tar -xzf laser.tgz
hexagon:r tkaiser$

hexagon:r tkaiser$ Rscript doinstall.R
. . .
. . .
hexagon:r tkaiser$ R CMD SHLIB mapit.c
```

A Prototype Model

- We will introduce a finite difference model that will serve to demonstrate what a computational scientist needs to do to take advantage of Distributed Memory computers using MPI
- The model we are using is a two dimensional solution to a model problem for Ocean Circulation, the Stommel Model

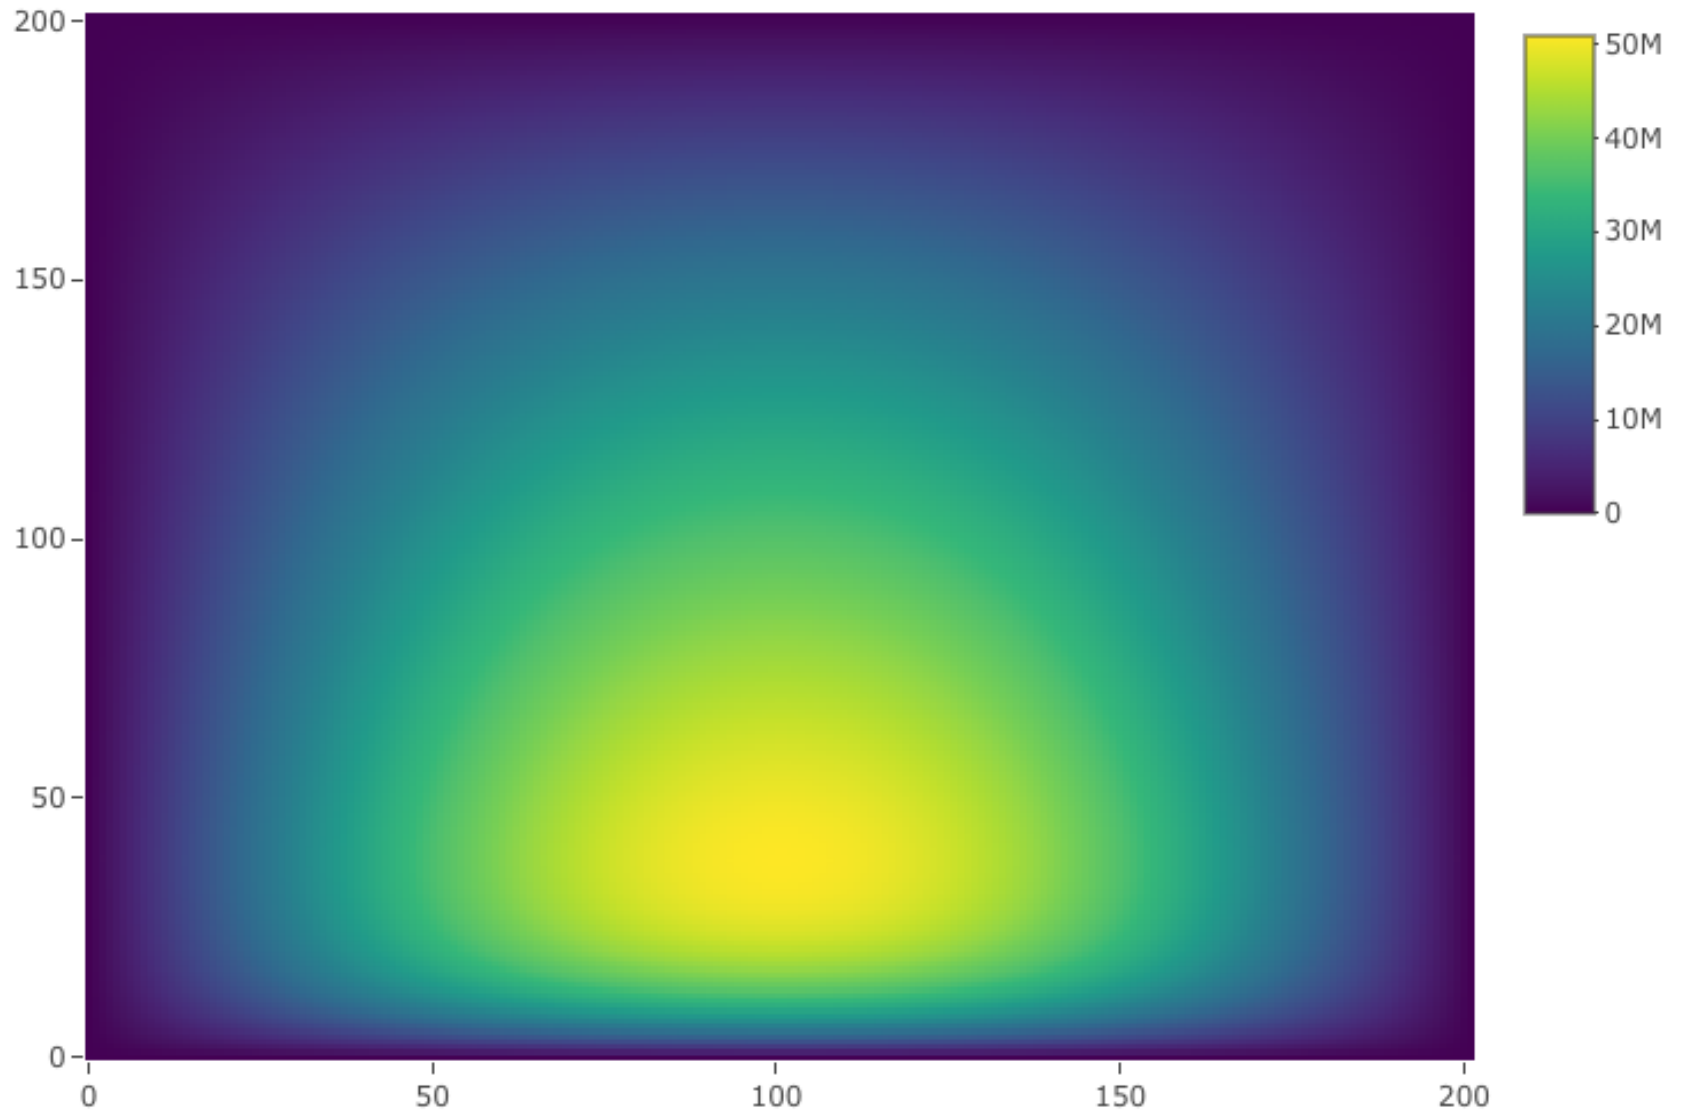
The Stommel Problem

- Wind-driven circulation in a homogeneous rectangular ocean under the influence of surface winds, linearized bottom friction, flat bottom and Coriolis force.
- Solution: intense crowding of streamlines towards the western boundary caused by the variation of the Coriolis parameter with latitude

Governing Equations Model

$$\begin{aligned} \psi &= 0 \\ \gamma \left(\frac{\partial^2 \psi}{\partial x^2} + \frac{\partial^2 \psi}{\partial y^2} \right) + \beta \frac{\partial \psi}{\partial x} &= f \\ f &= -\alpha \sin\left(\frac{\pi y}{2L_y}\right) \\ \psi &= 0 \end{aligned}$$
$$\begin{aligned} L_x &= L_y = 2000 \text{ Km} \\ \gamma &= 3 * 10^{(-6)} \\ \beta &= 2.25 * 10^{(-11)} \\ \alpha &= 10^{(-9)} \\ \psi &= 0 \end{aligned}$$

The steady state solution



R is **!!!!MUCH!!!!** slower

- Our “normal” example calculation 200x200 grid 75k steps
 - R > 2000 seconds
 - C/Fortran 3 seconds
- You can however, call C and Fortran compiled subroutines from R
- You can mix R and C/Fortran in MPMD fashion
 - MPI tasks 0 to n-2 might be Fortran or C
 - MPI task n-1 could be a R graphics program
 - Have example
 - `mpiexec -n 3 ./ccalc : -n 1 plot.R < small.in`

Domain Discretization

Define a grid consisting of points (x_i, y_j) given by

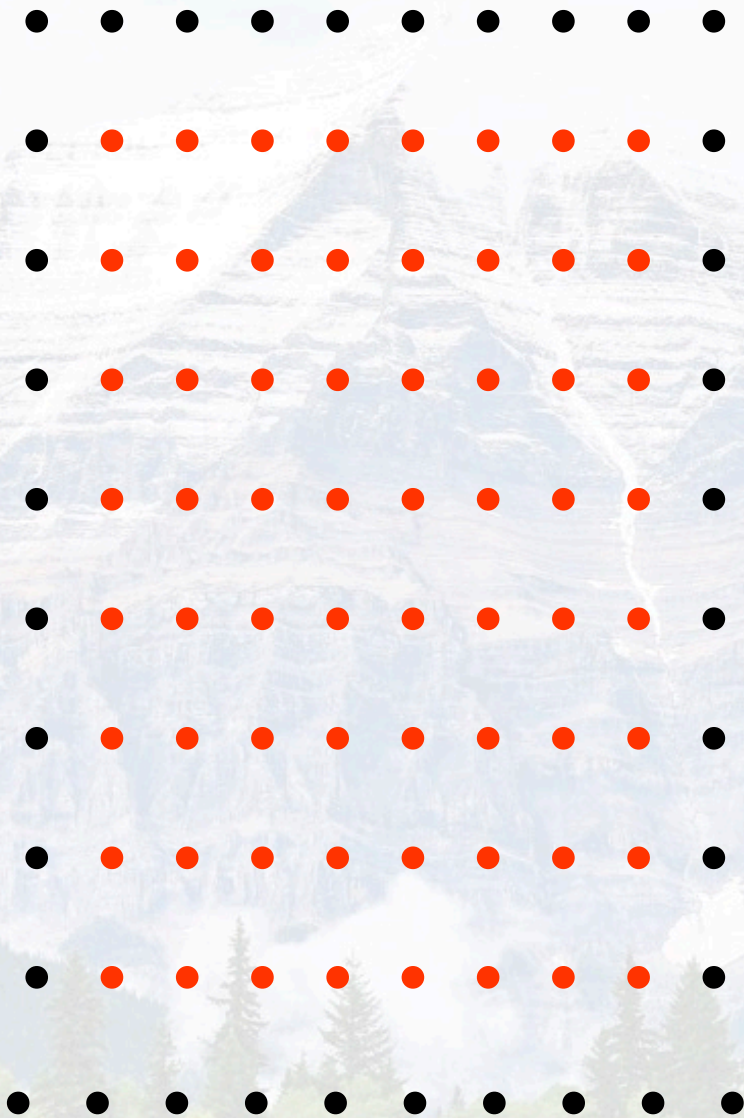
$$x_i = i\Delta x, i = 0, 1, \dots, nx+1$$

$$y_j = j\Delta y, j = 0, 1, \dots, ny+1$$

$$\Delta x = L_x / (nx + 1)$$

$$\Delta y = L_y / (ny + 1)$$

Domain Discretization



Seek to find an
approximate solution

$\psi(x_i, y_j)$ at points (x_i, y_j) :
 $\psi_{i,j} \approx \psi(x_i, y_j)$

Centered Finite Difference Scheme for the Derivative Operators

$$\frac{\partial \psi}{\partial x} \approx \frac{\psi_{i+1,j} - \psi_{i-1,j}}{2\Delta x}$$

$$\frac{\partial^2 \psi}{\partial x^2} \approx \frac{\psi_{i+1,j} - 2\psi_{i,j} + \psi_{i-1,j}}{(\Delta x)^2}$$

$$\frac{\partial^2 \psi}{\partial y^2} \approx \frac{\psi_{i,j+1} - 2\psi_{i,j} + \psi_{i,j-1}}{(\Delta y)^2}$$

Governing Equation Finite Difference Form

$$\psi_{i,j} = a_1\psi_{i+1,j} + a_2\psi_{i-1,j} + a_3\psi_{i,j+1} + a_4\psi_{i,j-1} - a_5f_{i,j}$$

$$a_1 = \frac{\Delta y^2}{2(\Delta x^2 + \Delta y^2)} + \frac{\beta \Delta x^2 \Delta y^2}{4\gamma \Delta x (\Delta x^2 + \Delta y^2)}$$

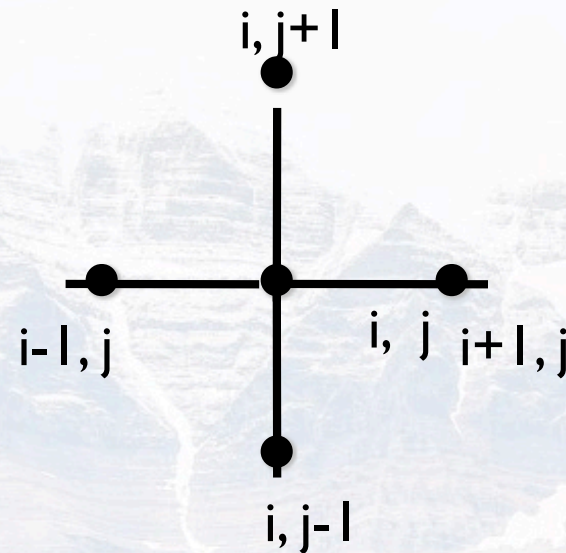
$$a_2 = \frac{\Delta y^2}{2(\Delta x^2 + \Delta y^2)} - \frac{\beta \Delta x^2 \Delta y^2}{4\gamma \Delta x (\Delta x^2 + \Delta y^2)}$$

$$a_3 = \frac{\Delta x^2}{2(\Delta x^2 + \Delta y^2)}$$

$$a_4 = \frac{\Delta x^2}{2(\Delta x^2 + \Delta y^2)}$$

$$a_5 = \frac{\Delta x^2 \Delta y^2}{2\gamma (\Delta x^2 + \Delta y^2)}$$

Five-point Stencil Approximation



interior grid points: $i=1, nx;$
 $j=1, ny$

boundary points:

$(i, 0) \text{ \& } (i, ny+1) ; i=0, nx+1$

$(0, j) \text{ \& } (nx+1, j) ; j=0, ny+1$

$$\psi_{i,j} = a_1 \psi_{i+1,j} + a_2 \psi_{i-1,j} + a_3 \psi_{i,j+1} + a_4 \psi_{i,j-1} - a_5 f_{i,j}$$

$$\psi_{i,0} = \psi_{i,ny+1} = 0; \quad \psi_{0,j} = \psi_{nx+1,j} = 0;$$

Jacobi Iteration

Start with an initial guess for $(\psi_{i,j})$
Set forcing function and Boundary Conditions

Repeat the process

do $i = 1, nx; j = 1, ny$

$$(\psi_{i,j})_{new} = a_1(\psi_{i+1,j}) + a_2(\psi_{i-1,j}) + \\ a_3(\psi_{i,j+1}) + a_4(\psi_{i,j-1}) - \\ a_5 f_{i,j}$$

end do

$$(\psi_{i,j}) = (\psi_{i,j})_{new}$$

R Version - Forcing Function

```
do_force<-function(i1,i2,j1,j2,alpha){  
  forf<-matrix(0.0,nrow=((i2-i1)+3),ncol=((j2-j1)+3))  
  r1<-range(0,(i2-i1)+2)  
  r2<-range(0,(j2-j1)+1)  
  for (i in r1[1]:r1[2]){  
    for (j in r2[1]:r2[2]){  
      y<-(j+j1-1)*dy  
      forf[i+1,j+1]<-(-alpha*sin(y*a6))  
    }  
  }  
  return(forf)  
}
```

R Version - Boundary Conditions

```
bc<-function(psi,i1,i2,j1,j2,nx,ny){  
  myid <-0  
  numprocs <- 1  
  myleft<-myid-1  
  myright<-myid+1  
  if(myleft <= -1){myleft<-(-1)}  
  if(myright >= numprocs){myright=(-1)}  
  if(myleft == -1){  
    psi[,1]=0.0  
  }  
  if(myright == -1){  
    psi[,ncol(psi)]=0.0  
  }  
  psi[1,]=0.0  
  psi[nrow(psi),]=0.0  
  return(psi)  
}
```


R Version - Set Up

```
ngrid<-scan("st.in",integer(),2)
grid<-scan("st.in",double(),2,skip=1)
const<-scan("st.in",double(),3,skip=2)
steps<-scan("st.in",integer(),1,skip=3)
nx<-ngrid[1]
ny<-ngrid[2]
lx<-grid[1]
ly<-grid[2]
alpha<-const[1]
beta<-const[2]
gamma<-const[3]
```

```
numprocs<-1
myid<-0
i1org<-1
i2org<-nx
j1org<-1
j2org<-ny
i1<-1
i2<-nx
j1<-1
j2<-ny
dj<-as.double(j2)/as.double(numprocs)
j1<-round(1.0+myid*dj)
j2<-round(1.0+(myid+1)*dj)-1
```

R Version - Calculation Function

```
do_jacobi<-function(psi){  
  # does a single Jacobi iteration step  
  # input is the grid and the indices for the interior cells  
  # new_psi is temp storage for the the updated grid  
  # output is the updated grid in psi and diff which is  
  # the sum of the differences between the old and new grids  
  gdiff<-0.0  
  is=2  
  ie=dim(psi)[1]  
  js=2  
  je=dim(psi)[2]  
  new_psi[,je]<-psi[,je]  
  new_psi[,1]<-psi[,1]  
  new_psi[ie,]<-psi[ie,]  
  new_psi[1,]<-psi[1,]  
  je=je-1  
  ie=ie-1  
  for (i in is:ie){  
    for (j in js:je){  
      new_psi[i,j]<-a1*psi[i+1,j] + a2*psi[i-1,j] + a3*psi[i,j+1] +  
a4*psi[i,j-1] - a5*forf[i,j]  
      gdiff<-gdiff+abs(new_psi[i,j]-psi[i,j])  
    }  
  }  
  return(new_psi)  
}
```


R Version - Main Loop

```
iout=steps/100
if(iout == 0){iout=1}
tymer(reset=T)
for(i in 1:steps){
  psi<-do_jacobi(psi)
  if(myid == 0){
    if (((i) %% iout) == 0){
      print(paste(i,gdiff,tymer()))
    }
  }
}

writeBin(as.vector(psi),"stom00.out")
```

MPI version of the Serial Code

Timothy H. Kaiser, Ph.D.

Overview

We will choose one of the two dimensions and subdivide the domain to allow the distribution of the work across a group of distributed memory processors

We will focus on the principles and techniques used to do the MPI work in the model

We will discuss the Rmpi version of the program there are C, PYthon, and Fortran versions also

Domain Decomposition (1d)

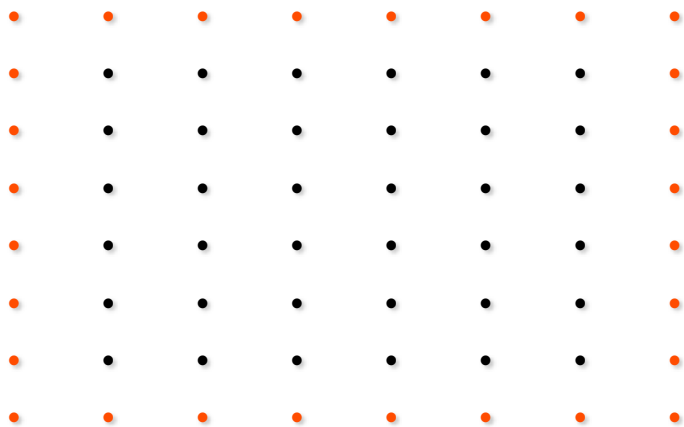
- We set our array bounds on each processor so that:
 - We take our original grid and break it into numnodes subsections of size $\text{nx}/\text{numnodes}$
 - Each processor calculates for a different subsection of the grid
 - No two processors calculate ψ for the same (i,j)
- We add special boundary cells for each subsection of the grid called ghost cells

Domain Decomposition (1d)

Physical domain is sliced into sets of columns so that computation in each set of columns will be handled by different

Serial Version

all cells on one processor

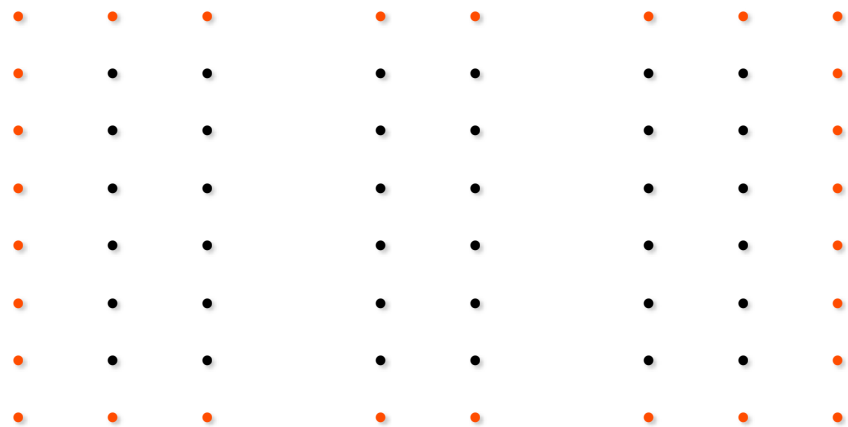


Parallel Version

node 0

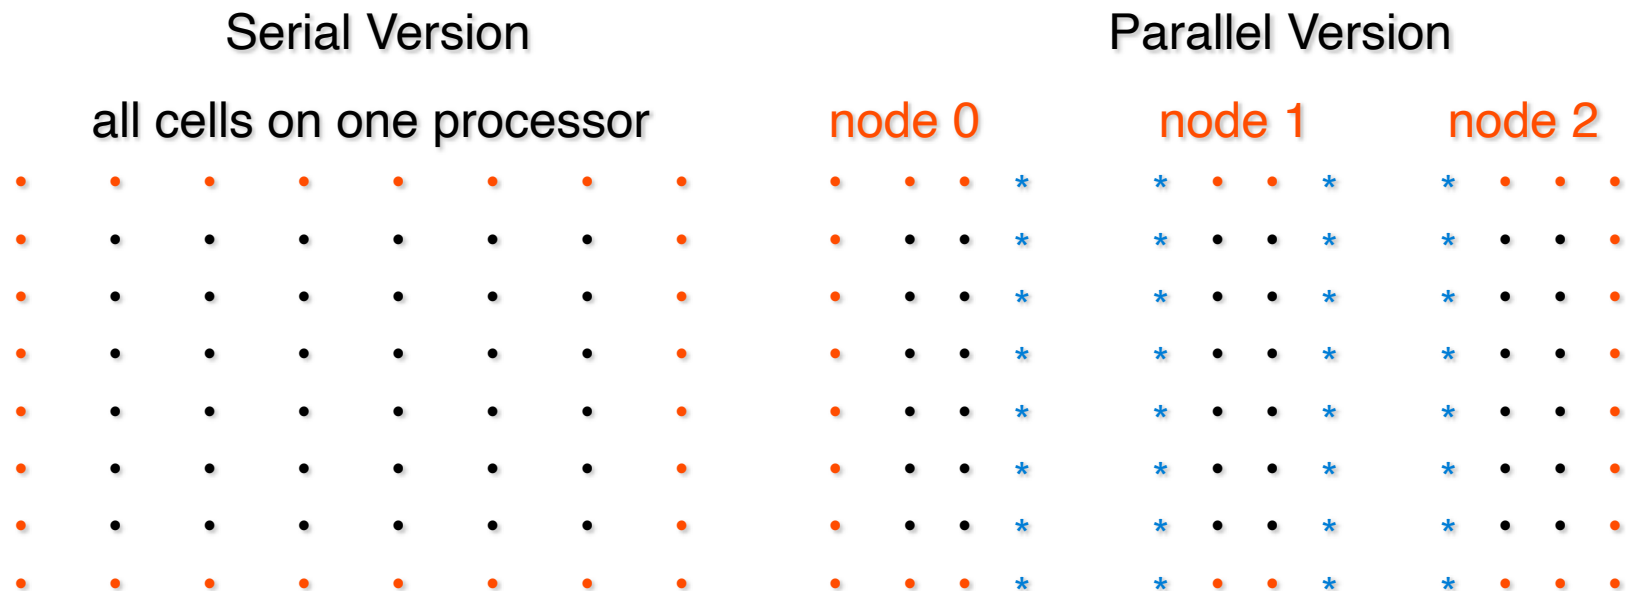
node 1

node 2



Domain Decomposition (1d)

With ghost cells our decomposition becomes...



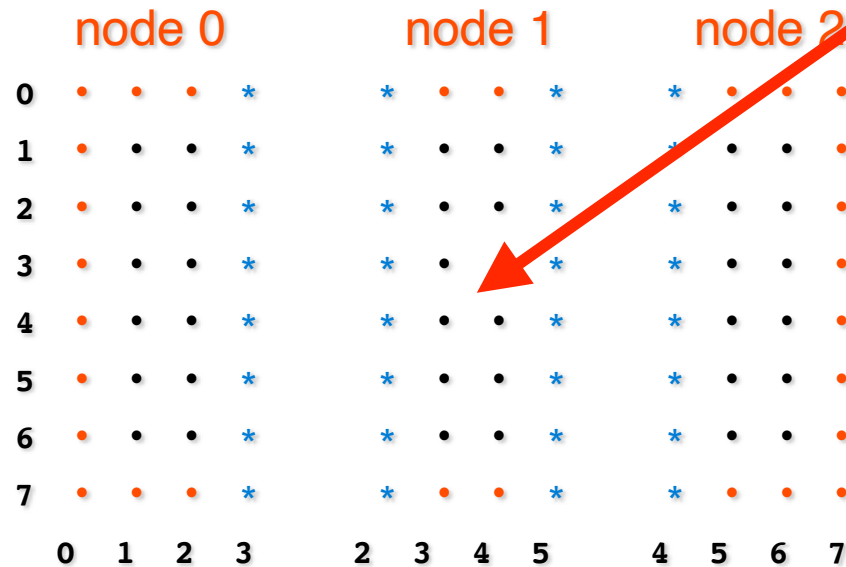
Domain Decomposition (1d)

How and why are ghost cells used?

Node 0 allocates space for $\psi(0:7,0:3)$ but calculates $\psi(1:6,1,2)$

Node 1 allocates space for $\psi(0:7,2:5)$ but calculates $\psi(1:6,3,4)$

Node 2 allocates space for $\psi(0:7,4:7)$ but calculates $\psi(1:6,5,6)$



To calculate the value for $\psi(4,4)$
node1 requires the value from
 $\psi(4,3), \psi(5,4), \psi(3,4), \psi(4,5)$

Where does it get the value for
 $\psi(4,5)$? From node2, and it holds
the value in a ghost cell

Ghost cell updates

- When do we update ghost cells?
- Each trip through our main loop we call `do_transfer` to update the ghost cells
- How do we do it?

How do we update ghost cells?

Processors send and receive values to and from neighbors

Need to exchange with left and right neighbors except processors on far left and right only transfer in 1 direction

Trick 1 to avoid deadlock:

Even # processors

send left

receive from left

send right

receive from right

Odd # processors

receive from right

send to right

receive for left

send to left

Trick 2 to handle the end processors

Send to MPI_PROC_NULL instead of a real processor

How do we update ghost cells?

It's a 4-stage operation

Example with 4 nodes:

	Proc 0	Proc 1	Proc 2	Proc 3
Stage 1	Send left to MPI_PROC_NULL	Receive right from Proc 2	Send left to Proc 1	Receive right from MPI_PROC_NULL
Stage 2	Receive left from MPI_PROC_NULL	Send right to Proc 2	Receive left from Proc 1	Send right to MPI_PROC_NULL
Stage 3	Receive right from Proc 1	Send right to Proc 0	Receive right from Proc 3	Send right to Proc 2
Stage 4	Send right to Proc 1	Receive right from Proc 0	Send right to Proc 3	Receive right from Proc 2

do_transfer

```
do_transfer<-function(psi,i1,i2,j1,j2){  
  mystat=0  
  myid <- mpi.comm.rank(comm=mpi_comm_world)  
  myleft<-myid-1  
  myright<-myid+1  
  if(myleft <= -1){myleft<-(-1)}  
  if(myright >= numprocs){myright=(-1)}  
  lc=dim(psi)[2]  
  if(is.even(myid)){  
    ...  
  }else{  
    ...  
  }  
  return(psi)  
}
```

STEP1 & 2: introduce the MPI environment and read data

```
mpi_comm_world<-0
myid <- mpi.comm.rank(comm=mpi_comm_world)
numprocs <- mpi.comm.size(comm=mpi_comm_world)
myname <- mpi.get.processor.name()

if (myid != 0) {
  nx<-0
  ny<-0
  lx<-0.0
  ly<-0.0
  alpha<-0.0
  beta<-0.0
  gamma<-0.0
  steps=0
}else{
  ngrid<-scan("st.in",integer(),2)
  grid<-scan("st.in",double(),2,skip=1)
  const<-scan("st.in",double(),3,skip=2)
  steps<-scan("st.in",integer(),1,skip=3)
  #bcast should be here
  nx<-ngrid[1]
  ny<-ngrid[2]
  lx<-grid[1]
  ly<-grid[2]
  alpha<-const[1]
  beta<-const[2]
  gamma<-const[3]
}
```


STEP 3: bcast the data

```
## mpi.bcast(nx,      type=1,comm = mpi_comm_world)
nx<-mpi.bcast(nx,      type=1,comm = mpi_comm_world)
## mpi.bcast(ny,      type=1,comm = mpi_comm_world)
ny<-mpi.bcast(ny,      type=1,comm = mpi_comm_world)
##      mpi.bcast(steps, type=1,comm = mpi_comm_world)
steps<-mpi.bcast(steps, type=1,comm = mpi_comm_world)
lx<-mpi.bcast(lx,      type=2,comm = mpi_comm_world)
ly<-mpi.bcast(ly,      type=2,comm = mpi_comm_world)
alpha<-mpi.bcast(alpha, type=2,comm = mpi_comm_world)
beta<-mpi.bcast(beta,  type=2,comm = mpi_comm_world)
gamma<-mpi.bcast(gamma, type=2,comm = mpi_comm_world)
print(paste(myid,nx,ny,steps,lx,ly,alpha,beta,gamma))
```

STEP4 : Set the constants

```
i1org<-1
i2org<-nx
j1org<-1
j2org<-ny
i1<-1
i2<-nx
j1<-1
j2<-ny
dj<-as.double(j2)/as.double(numprocs)
j1<-round(1.0+myid*dj)
j2<-round(1.0+(myid+1)*dj)-1

dx<-lx/(nx+1.0)
dy<-ly/(ny+1.0)
dx<-lx/(nx+1.0)
dx2<-dx*dx
dy2<-dy*dy
bottom<-2.0*(dx2+dy2)
a1<-((dy2/bottom)+(beta*dx2*dy2)/(2.0*gamma*dx*bottom))
a2<-((dy2/bottom)-(beta*dx2*dy2)/(2.0*gamma*dx*bottom))
a3<-dx2/bottom
a4<-dx2/bottom
a5<-dx2*dy2/(gamma*bottom)
a6<-pi/(ly)
gdiff<-0.0
```


Step 5: Set up

```
psi<-matrix(1.0,nrow=((i2-i1)+3),ncol=((j2-j1)+3))  
  
print(paste(myid,"covers",i1,i2,j1,j2))  
  
psi<-bc(psi)  
  
new_psi<-matrix(0.0,nrow=((i2-i1)+3),ncol=((j2-j1)+3))
```

```
bc<-function(psi,i1,i2,j1,j2,nx,ny){  
  myid <-0  
  numprocs <- 1  
  myleft<-myid-1  
  myright<-myid+1  
  if(myleft <= -1){myleft<-(-1)}  
  if(myright >= numprocs){myright=(-1)}  
  if(myleft == -1){  
    psi[,1]=0.0  
  
  }  
  if(myright == -1){  
    psi[,ncol(psi)]=0.0  
  }  
  psi[1,]=0.0  
  psi[nrow(psi),]=0.0  
  return(psi)  
}
```

```
bc<-function(psi,i1,i2,j1,j2,nx,ny){  
  myid <- mpi.comm.rank(comm=mpi_comm_world)  
  numprocs <- mpi.comm.size(comm=mpi_comm_world)  
  myleft<-myid-1  
  myright<-myid+1  
  if(myleft <= -1){myleft<-(-1)}  
  if(myright >= numprocs){myright=(-1)}  
  if(myleft == -1){  
    psi[,1]=0.0  
  
  }  
  if(myright == -1){  
    psi[,ncol(psi)]=0.0  
  }  
  psi[1,]=0.0  
  psi[nrow(psi),]=0.0  
  return(psi)  
}
```

Residual

- In our serial program, the routine `do_jacobi` calculates a residual for each iteration
- The residual is the sum of changes to the grid for a jacobi iteration
- Now the calculation is spread across all processors
- To get the global residual, we can use the `MPI_Reduce` function

Final Results

- In our parallel program, the grid is spread across all processors
- Each processor can print its portion of the grid
- Regular MPI has special parallel IO routines (difficult)
- Here we "just" bring all of the data back to process 0 and save it

Our Main Loop

```
iout=steps/100
if(iout == 0){iout=1}
if(myid == 0){tymer(reset=T)}
for(i in 1:steps){
  psi<-do_transfer(psi,i1,i2,j1,j2)
  psi<-do_jacobi(psi,i1,i2,j1,j2)
  mydiff<-gdiff
  todiff<-mpi.reduce(mydiff, type=2, op="sum",dest = 0, comm = mpi_comm_world)
  if (((i) %% iout) == 0){
    print(paste(myid,max(psi),min(psi)))
    if(myid == 0){
      print(paste(i,todiff,tymer()))
    }
  }
}
getresults(psi,nx,ny)

bonk<-mpi.finalize()
```

do_transfer

```
do_transfer<-function(psi,i1,i2,j1,j2){  
  mystat=0  
  myid <- mpi.comm.rank(comm=mpi_comm_world)  
  myleft<-myid-1  
  myright<-myid+1  
  if(myleft <= -1){myleft<-(-1)}  
  if(myright >= numprocs){myright=(-1)}  
  lc=dim(psi)[2]  
  if(is.even(myid)){  
    ...  
  }else{  
    ...  
  }  
  return(psi)  
}
```

do_transfer

```
    if(is.even(myid)){  
# we are on an even col processor  
# send to left  
    if(myleft != -1){  
        #print(paste(myid,"to",myleft))  
        mpi.send(psi[,2], type=2, dest=myleft, tag=100, comm = mpi_comm_world)  
# rec from left  
        #print(paste(myid,"from",myleft))  
        psi[,1]<-mpi.recv(psi[,1], type=2,source=myleft,tag=100,comm = mpi_comm_world,status = mystat)  
    }  
# rec from right  
    if(myright != -1){  
        #print(paste(myid,"from",myright))  
        psi[,lc]<-mpi.recv(psi[,lc],type=2,source=myright, tag=100,comm = mpi_comm_world,status = mystat)  
# send to right  
        #print(paste(myid,"to",myright))  
        mpi.send(psi[,lc-1], type=2, dest=myright, tag=100, comm = mpi_comm_world)  
    }  
}
```


do_transfer

```
else{
# we are on an odd col processor
# rec from right
    if(myright != -1){
        psi[,lc]<-mpi.recv(psi[,lc],type=2,source=myright,tag=100,comm = mpi_comm_world, status = mystat)
# send to right
        mpi.send(psi[,lc-1],type=2,dest=myright,tag=100,comm = mpi_comm_world)
    }
# send to left
    if(myleft != -1){
        mpi.send(psi[,2], type=2, dest=myleft, tag=100, comm = mpi_comm_world)
# rec from left
        psi[,1]<-mpi.recv(psi[,1],type=2,source=myleft,tag=100,comm = mpi_comm_world, status = mystat)
    }
}
```

getresults

```
getresults<-function(psi,nx,ny){  
  myid <- mpi.comm.rank(comm=mpi_comm_world)  
  tosend=(nrow(psi)-2)*(ncol(psi)-2)  
  rcounts<-vector("integer",numprocs)  
  rcounts<-mpi.gather(tosend, 1, rcounts, root = 0, comm = mpi_comm_world)  
  typeof(rcounts)  
  typeof(psi)  
  tot<-sum(rcounts)  
  full<-vector("double",tot)  
  full<-mpi.gatherv(psi[(2:(nrow(psi)-1)),(2:(ncol(psi)-1))],2, full,  
                    rcounts, root = 0, comm = mpi_comm_world)  
  if(myid == 0){writeBin(full,"stom01.out")}  
}
```

```
#install.packages("plotly", lib="/Library/Frameworks/R.framework/Versions/3.6/Resources/library")  
#result<-readBin("stom01.out",double(),10000)  
#grid<-matrix(result,100,100)  
#library(plotly)  
#plot_ly(z =grid)
```